

Abhinav Kotta

(689) 209-8517 | abhinavkotta.io@gmail.com | <https://www.linkedin.com/in/akotta03/> | <https://abhinavkotta.com/> | <https://github.com/abhinav-kotta/>

EDUCATION

US CITIZEN

Cornell University - Cornell Tech

Master of Engineering (M.Eng.), Computer Science

Aug 2025 – May 2026

NYC, NY

University of Central Florida

Bachelor of Science (B.S.), Computer Science

Aug 2021 – May 2025

Orlando, FL

PROFESSIONAL EXPERIENCE

L3Harris Technologies

Software Engineer

Jul 2026 – Present

Melbourne, FL

- Architected a situational-awareness system for **SD-WAN** to predict tunnel degradation.
- Built a recommendation engine using **ELK** stack, weighted health scoring, and deterministic remediation logic.
- Evaluated **Isolation Forest** and **HMMs** to replace rule-based scoring with ML-driven anomaly detection.
- Built anomaly-detection prototypes to identify abnormal tunnel behavior early.

TheMindOverMarket

CTO

Jan 2026 – May 2026

NYC, NY

- Architected a deterministic trading-rule engine with finite primitives and real-time behavioral enforcement
- Built a SaaS strategy supervisor for intraday traders and prop firms enforcing execution discipline.
- Selected as a **Top 11 finalist** at Cornell Tech Startup Awards for behavioral trading platform innovation
- Integrated broker APIs to detect, quantify, and penalize deviations from predefined trading plans in real time.

Los Alamos National Laboratory

Software Engineer Intern

May 2025 – Aug 2025

Los Alamos, NM

- Built a platform that simplified HPC access for physicists using **Slurm**, cutting job submission time by **50%**
- Configured a 10-node **HPC** cluster with distributed **GPUs** and Infiniband for high-throughput scientific computing
- Deployed **GPU-enabled Docker** containers with **Charliecloud** and **Docker** for secure HPC execution
- Improved resource efficiency by providing real-time **GPU/CPU** using **Grafana** to reduce underutilization

L3Harris Technologies

Software Engineer Intern

May 2024 – Aug 2024

Melbourne, FL

- Developed an **AI-powered** subject matter expert to preserve critical business knowledge
- Containerized **AWQ-quantized LLMs** with **Docker**, boosting inference accuracy by **40%**
- Engineered high-throughput **microservices backend** powering production-grade **RAG** with fine-tuned LLMs
- Deployed resilient inference systems with **Kubernetes**, scaling AI/ML workloads and improving throughput by **50%**

PROJECTS AND LEADERSHIP

SMB Finance Ops | 2026 AI Agent Hackathon @ Datadog | [Github](#)

Feb 2026

- Built event-driven FinanceOps agent automating collections, reconciliation, and bill pay from call and payment data.
- Integrated Velma Transcribe for commitment extraction and enforced safe payments via constrained delegation.
- Implemented RL with Agent Lightning to optimize follow-up and escalation from real cashflow outcomes.

AI Market Analytics | 2025 Hacklytics @ Georgia Tech | [Github](#)

Feb 2025

- Built a **Typescript/React** app using **Node** and **FastAPI** to generate and analyze AI-driven sales conversations
- Implemented message communication between **LLMs** using Google Gemini and AWS Polly
- Deployed a dynamic web application on **AWS EC2**, configured with **Caddy** for domain-based routing

TECHNICAL SKILLS

Programming & Scripting: OOP, Python, C++, Java, C, C#, JavaScript, TypeScript, SQL, Bash, HTML

Data & AI Technologies: Pandas, Claude, Machine Learning, PyTorch, RAG, Generative AI, Inference, PostgreSQL

Systems & Infrastructure: Docker, Kubernetes, JUnit, HPC, Slurm, Infiniband, Ethernet, GPU, CUDA, Git, CLI

Backend & Cloud: AWS (Lambda, S3, API Gateway), MongoDB, Docker, Git